
AOS Group 3 - SOS Design Documentation

19th October, 2015

Contents

1	Overview	3
2	Architecture	3
2.1	Execution Model	3
2.2	Events	4
2.3	System Calls	5
3	System Memory	6
3.1	Frame Table	6
3.2	Swap File	7
4	Processes and Virtual Memory	8
4.1	Process Management	8
4.2	Page Table	8
4.3	Page Replacement	9
4.4	Regions	11
4.5	Page Faults	11
5	I/O and Drivers	12
5.1	IO Vectors	12
5.2	I/O Device Abstraction	13
5.3	Serial Device	14
5.4	NFS	14
5.5	Timer Interrupts	15

1 Overview

The Simple Operating System, or SOS, performs resource management for applications running on top of the seL4 Microkernel on ARMv7, specifically the SABRE Lite i.MX6. This document describes the design of SOS as implemented by AOS Group 3. For each key design choice in the project, it also discusses benefits and trade-offs of the decision, and possible alternative approaches. Refer to Doxygen documentation for implementation-specific details and conventions.

There are 4 distinct problem-areas that have been explored in SOS: the architecture and execution model, system memory and swapping, processes and virtual addresses and I/O and drivers. Understanding the architecture of SOS, the event loop and syscall dispatch are essential for understanding subsequent discussions, and thus offers a solid starting point for introducing SOS.

2 Architecture

2.1 Execution Model

2.1.1 Design

Implemented in: `/apps/sos/src/main.c`

At its core, SOS uses an event-based model implemented using *continuations*. A continuation allows for non-blocking behaviour during otherwise by blocking in a single-threaded implementation (for example I/O) by providing a way to preserve execution state, thereby maximising the time where SOS is able to accept new client requests. This allows for a single-threaded implementation, while preserving many of the performance characteristics of a multi-threaded implementation.

A continuation is used to track the current state of execution, including for example, the original syscall number or fault address and the number of bytes read or written during the current entry to the kernel. A continuation is unique to each *process*. This model allows execution to resume once an interrupt occurs, and thus permits SOS non-blocking I/O.

2.1.2 Discussion

A multi-threaded kernel does offer performance advantages on SMP-capable hardware through the virtue of being able to allow entry from multiple clients simultaneously, the benefits of multi-threading in the environment SOS is designed for are somewhat diminished. Foremost, on ARM, seL4 supports only a single core and thus a single-threaded implementation may offer improved performance over a threaded implementation due to the costs imposed by locking, though also there is only limited difference under non-high-performance workloads due to the prevalence of locking, which limits scalability in multi-threaded systems.

The implementation of continuations in SOS is conceptually simple, and allows for type-checking at compile-time, which affords some correctness guarantees. However, this struct becomes exceedingly large as more behaviours are added to SOS, each of which may depend on and need to extend upon the state already stored in the continuation.

Yet more sophisticated data structures could also be used, which are amenable to both static analysis or better scalability, such as a hash-table, though to achieve both static analysis and scalability requires a yet more sophisticated data structure which imposes added complexities outweighing the benefits offered to the current incarnation of SOS.

2.2 Events

Implemented in: `/libs/libclock/src/main.c`

2.2.1 Design

The `event_loop()` serves as the central point of dispatch within SOS, and services three different sources of events:

1. Ready queue
2. Interrupts
3. System calls

Network and Timer interrupts may come over Asynchronous IPC. The behaviour for these will be described in their respective sections later in the document. System Calls are received from processes, and will be discussed in depth in the following section.

Processes in the Ready queue are serviced before either Interrupts or System Calls. Processes are enqueued by interrupt handlers: once an interrupt request on behalf of the process is complete, that process is queued for selection in the event loop. Conversely,

the current flow of execution can indicate that it is waiting on an interrupt through the use of `longjmp` to `ipc_event_buf`, re-entering the event loop.

At the completion of a client request, the function `syscall_end_continuation()` is used to reply to the client in a consistent manner, ensuring that all execution state is released and reset appropriately. The caller to `syscall_end_continuation` indicates the process to reply to, the return code, and whether or not the request was successful. At the completion of a request, the state of the process' continuation is reset.

2.2.2 Discussion

Earlier implementations of the process 'ready queue' a stack, though this proved prone to starvation: earlier processes may never be serviced. A queue is now used to guarantee fairness.

2.3 System Calls

Implemented in: `/apps/sos/src/handler.c`, `/apps/sos/src/syscall.c`

2.3.1 Design

To facilitate re-entry of the syscall handlers, each system call is implemented in two parts: a setup handler, and an execution handler. The setup handler is called once, at the beginning of a syscall and before the execution handler is invoked. This allows a single, consistent means of storing any arguments provided via IPC within the continuation, so that they may be available to the execution handler.

The syscall dispatch is represented as a table which is indexed by the system call number and the handler type, giving simple and $O(1)$ access to the syscalls implemented within SOS.

During the execution of a system call, two concepts are used to access the current current state:

current_process The process which invoked the current system call.

effective_process The process which SOS is acting on behalf of. This is usually identical to `current_process`, with the exception being spawning a child process. It is defined in terms of the current process.

2.3.2 Discussion

The approach to splitting the setup and execution handler for the syscalls proved largely successful and simplified the implementation substantially. However, within the execution handler, the approach to continuations whereby numerous restartable code-blocks may be contained in a single function, could prove challenging to debug.

Extending this concept across the entirety of the continuation, whereby a continuation would contain a list of outstanding execution events which may (or may not) be required to satisfy each syscall could be a viable approach and simplify both the syscall implementation, and the continuation structure. Currently many fields in the continuation denote execution state, and this approach would allow that state of execution to be separated from other state and represented consistently.

The downside of this approach would be that the linked list largely becomes a complex representation of a function call.

3 System Memory

3.1 Frame Table

Implemented in: `/apps/sos/src/frametable.c`

3.1.1 Design

The frame table is the representation of physical memory. This memory is mapped by applications and also used by SOS. At boot-time the frame table is initialised for the entirety of memory, with frames backing the frame table being allocated at its beginning. The frame table only supports `seL4_SmallPageObjects`.

The frame table is an array of structs, where each struct contains the details specific to the frame (namely, its address and `CPtr`), as well as a pointer to the next element.

This structure contains as a linked-list of free frames, which gives $O(1)$ performance for allocation and de-allocation of frames. Similarly, as the underlying structure is an array, it also offers $O(1)$ lookup as the address can be transformed in to an index into the array trivially.

3.1.2 Discussion

As the data-structure is 12 bytes in size under the target architecture, this results in approximately 768 frames of memory being consumed per GiB of physical memory, delaying boot time.

Furthermore, as the data-structure for a frame doubles to provide pointers to the next free element, under high memory contention situations approximately 1/3 of the frame table size will be used to store NULL pointers. Improving the representation of the frame table would allow for another 256 frames to be made available for use per GiB of memory. However, this equates to 1MiB of memory per GiB, which is not significant for most applications.

3.2 Swap File

Implemented in: `/apps/sos/src/swap.c`

3.2.1 Design

When the system encounters contention for the available physical memory resources, some memory may be swapped to disk. The data structure which underpins the swap file is similar to that which underpins the frame table: an array whose free nodes form a linked list, and as such also offers constant time lookup, allocation and deallocation.

When reading data from the swap file, or otherwise freeing it, the swap file is not altered.

The swap structure also features a check-sum field to protect against altered or corrupted data from being read-in from disk.

3.2.2 Discussion

As the swap file is very similar in design to the frame table, it has many of the same drawbacks. The structure representing a swap file uses 8 byte nodes, yielding 512 frames consumed per GB of swap.

While pre-allocating the frame table could be argued for: the memory it consumes would be consumed anyway, and it trades off boot time performance for run-time performance, this argument cannot be made with the swap file. Indeed, the system may be forced to perform swapping due to the frames occupied by the data-structure representing the swap file.

A preferable solution would be the use of a data structure which allocated memory in frames on demand, which could be represented as a two-level ‘page table’. However, such an approach would fix the size of the swap file which may be a yet more significant issue on systems with restricted disk space or where the size of the swap file should otherwise be restricted.

4 Processes and Virtual Memory

4.1 Process Management

Implemented in: `/apps/sos/src/process.c`, `/apps/sos/src/elf.c`

4.1.1 Design

The design for Process IDs allows for $O(1)$ creation, lookup and removal, and will reuse process IDs minimally. This is achieved via an array of ‘PID entries’, where each entry in the array is a doubly-linked node, and can thus represent lists of both running and free PIDs. A PID entry contains a process ID and a flag to indicate whether it corresponds to a running process.

Closely related to the use of process IDs is the means by which a process is identified during a callback. SOS accounts for potential race conditions by using a reference to a ‘callback’ structure, which contains the creation time of the callback, and the PID corresponding to the process which the callback was created on behalf of. Given that the creation time of the process is known, it can be determined whether the process now owning the PID for the callback is the currently running process; a process starting after the creation time of the callback cannot own the callback, and thus the callback is discarded.

4.2 Page Table

Implemented in: `/apps/sos/src/addrspace.c`

4.2.1 Design

SOS implements a two-level page table within a process’ address space. The page table divides the virtual address as follows:



Figure 1: Address Format for Page Table lookup

This structure allows for $O(1)$ lookup: the upper 10 most bits index into the Page Directory, the next upper 10-most bits index in to the second level of the Page Table. This scheme allows each the Page Directory and the second level of the Page Table to address 1024 4-byte entries within one 4096-byte frame: each 4-byte entry is a pointer to a Page Table Entry.

A Page Table Entry (PTE) stores the frame address and page cap corresponding to the virtual address. As only 20 bits of the frame address are actually used, the lower 12 bits in the PTE are available for other purposes. Currently, 3 of these lower-most bits are used to provide flags indicating: whether the page data is stored on-disk, whether the page is pinned, and whether it has been referenced. (See ‘Page Replacement’ for details.)




Figure 2: Address Format as stored in the Page Table Entry

A linked list is used as the internal seL4 Page Directory and Page Tables (which are 12 and 8 bits, respectively) do not map cleanly in to the SOS Page Table design. The caps stored in this list exist only for the purpose of releasing the resource.

4.3 Page Replacement

Implemented in: `/apps/sos/src/page_replacement.c`

4.3.1 Design

SOS implements a second-chance local page replacement algorithm, which takes place over a number of steps:

1. Select the process from which to evict a page.
2. Select a page from that process using a second chance eviction model.
3. Write the page to the swap file.

Selecting the process to evict from is based on a threshold applied across all process: whether the process exceeds the average number of ‘evictable’ pages across all processes. The first process identified which satisfies this requirement is selected for eviction, or if

there isn't any process satisfying this requirement, the page is evicted from the current process.

In order to guarantee fairness of the system, the process for which eviction takes place is tracked, and subsequent selection takes place from the next process. Similarly, in order to prevent starvation as a result of two or more processes competing indefinitely for frames, so that they are able to make progress while the system is under very high resource constraints, the algorithm may randomly select the current process for eviction.

The page replacement algorithm uses a circular linked list of pages. For each page, if the referenced bit for that page table entry is found to be zero, the page is selected, otherwise the referenced bit is set to zero and the algorithm proceeds to the next allocated page.

Pages which are pinned, or already swapped to disk are not considered valid for selection. Thus there is the potential that there are no pages in the process which are valid selection. In this event, (i.e., the event we iterate through the list more than twice), the process which SOS is acting on behalf of is killed in order to relieve the memory contention.

Once a page has been selected, the contents of the frame are written to the next available location in the swap file. During the copy, the PTE is marked as **pinned**, and so is not valid for subsequent selection. Finally once the write is complete the page is marked as **swapped**, indicating an address on disk.

4.3.2 Discussion

Designing a scheme which maximised the amount of memory available for eviction, thereby maximising the number of processes available to run on the system, while preserving fairness received a significant amount of attention. While evicting from the currently running process, or even its parent, is naive approach does not make effective use of system resources.

The final implementation extends local replacement, and takes into account that if pages were to be evicted from any process, it must be done in such a way that two or more processes would not compete for the same resource indefinitely.

There are three factors to the process selection mechanism which help to ensure fairness:

queue-based selection once a process is assessed for selection, every other process on the system will be assessed before assessing that process again.

threshold based selection a process will only be selected should it have more than the average number of pages available for eviction

biased towards current__process selection is biased towards the current process (i.e., the process requesting pages) both by defaulting to this process should no other process satisfy the requirements for eviction, and by the random chance to select

this process, which serves to defeat processes endlessly competing for resources

As it is too expensive to consider all pages across all processes for eviction, it depends upon the correct accounting for the number of pages available to a process for eviction to calculate the threshold predicate, exposing the potential for future implementation mistakes.

4.4 Regions

Implemented in: `/apps/sos/src/addrspace.c`

4.4.1 Design

Regions are implemented as a singly-linked list. Each region also contains start, end, rights, and an offset for the starting location of that region in the elf file.

4.4.2 Discussion

There is no optimisation on the lookup cost of a region. This could be optimised through the use of binary search or similar approach, though is unnecessarily complex for most applications.

4.5 Page Faults

Implemented in: `/apps/sos/src/handler.c`

4.5.1 Design

Page faults within SOS occur in a number of circumstances:

- for lazy-loading data in from an ELF binary;
- faulting-in new pages within the heap or stack;
- when the process attempts to read data, which has been paged to disk, back into memory, and
- as a means of setting the frame referenced bit for second chance replacement

In the event a page corresponding to a virtual address has not been created within a processes address space when it faults, a new page will be created dynamically. When this occurs the corresponding region is used to determine whether any data should subsequently be loaded in to that page. This is indicated by an elf file offset being defined for the corresponding region. Therefore in this event content for code and data segments will be loaded from the ELF binary, while zeroed pages will be allocated for the stack and heap as the process faults within those regions.

Data paged to disk will be indicated by the **swapped** bit in the corresponding PTE. In the event a fault occurs on a page with this bit set, the address (which corresponds to an offset within the swap file) will be used to read in the page. During the read, another page may or may not be evicted depending on whether physical memory is available to facilitate the replacement.

The final scenario in which page faults will occur is in order to set the ‘referenced’ bit (**refd**) of the page to facilitate second-chance page replacement. During the replacement algorithm, pages where the referenced bit is unset will also be unmapped from the process’ address space. Thus when a fault occurs on a page where the referenced bit is not set, though it is still in memory, the referenced bit will be set and the page mapped back in to the process’ address space.

5 I/O and Drivers

5.1 IO Vectors

Implemented in: `/apps/sos/src/file.c`

5.1.1 Design

An IO Vector list is a representation of client or SOS-internal pages, allowing memory to be mapped for I/O on-demand.

An IO Vector provides the following fields:

vstart The starting address of the frame or virtual address (conditional)

sz The length of the operation to occur. **sz** will be set such that the operation does not exceed the page given the starting address **vstart**.

next The following IO Vector node.

sos_iov_flag Flag indicating whether the IO Vector node corresponds to a SOS frame, or process virtual address (thereby indicating whether address translation is required)

The benefit of the IO Vector list is that it allows for otherwise non-contiguous pages to be treated as though contiguous. It also provides a means for representing the progress of an IO operations, where the IO Vector list need simply be updated to reflect the current progress of an operation. Similarly, we need only guarantee the page for the current IO Vector is in memory, thus alleviating the need to read all pages for the current IO operation into memory.

IO Vectors are used to represent most I/O operations in SOS.

5.1.2 Discussion

Currently the usage of the IO Vector means that we may do short IO ops (e.g. 10 bytes to finish reading a page) considering only a single IO Vector at each time, even though we could do a large IO op (4k) and write across IO Vectors simultaneously. Correcting this would incur some complexity in the implementation, though also offers the opportunity for significant performance improvements, in particular when reading files over NFS.

5.2 I/O Device Abstraction

5.2.1 Design

The I/O Device is an abstraction over both block and character devices. In the case of SOS, this abstracts both Serial and NFS to create a uniform IO interface through mapping function pointers to their respective implementations.

Thus an IO Device provides an interface for all available IO operations over any driver. Thus it currently exposes:

- open
- close
- read
- write
- getdirent
- stat

For both the serial and NFS drivers. As the serial driver does not support some of these operations, those operations are undefined (NULL).

5.2.2 Discussion

The IO Device is a light-weight mechanism to abstract access to IO-related drivers. An alternative to this would be the implementation of a VFS-like abstraction. However, given that the goals of the project were to support a single, flat filesystem, and a serial device, this was deemed an overly extravagant abstraction.

5.3 Serial Device

Implemented in: `/apps/sos/src/serial.c`

5.3.1 Design

The serial device implementation satisfies single-reader / multiple-writer semantics, where the process is only permitted to read should no other process have the serial device open for reading. Reads in SOS are buffered, though only the most recent line is preserved. This provides bounds over the size of the buffer, while allowing non-blocking behaviour should have already been provided from the user.

When the `read()` is invoked by the client, the pages corresponding to that read are pinned, and subsequently unpinned when the read operation returns to the caller. As the buffer is at most 1024 bytes, up to two pages can be pinned because of this.

The serial driver is initialised at boot-time.

5.4 NFS

Implemented in: `/apps/sos/src/sos_nfs.c`

5.4.1 Design

The NFS implementation in SOS relies heavily on callbacks and continuations, which have been discussed previously in order to achieve non-blocking I/O operations within SOS. Prior to spawning a request for which a subsequent callback is expected, a ‘callback info’ struct will be allocated, which contains the time the callback was initiated and the process initiating the callback. The address of this is used as the token, and uniquely identifies the context for the callback once the callback is realised.

The PID from the callback is used to set the context for the subsequent execution, in the same way as the badge for a process' IPC is used to set the context for a syscall.

Detection of completion of the syscall is handled individually in each callback. In the event of error, or completion, the callback is responsible for triggering an appropriate reply to the client over IPC.

5.5 Timer Interrupts

Implemented in: `/libs/libclock/src/clock.c`

5.5.1 Design

The timer interrupt was implemented using the GPT timer on the SABRE.

The GPT timer allows multiple Output Compare Registers to be set simultaneously, and thus affords the ability to support two different types of interrupts.

1. Regular 100ms kernel tick
2. Timed callback events

Both timed callbacks and kernel tick callbacks are exposed via an interface for registering callback functions when these events occur.

Timed callbacks are called after a timeout expires. This is implemented by ensuring that all registrations are ordered at registration-time, and setting the corresponding Output Compare Register to the next event in the queue. The next event in the queue is determined through looking up the next callback in an array, which is sorted by callback time. Timer overflow is accounted for by searching for the event such that the difference between the current time, and the event time is minimal at the time of registration.

Timed callbacks are removed after being fired and are used to implement the `sleep` syscall. As kernel tick events satisfy different semantics to timed callback events in the implementation, they are distinguished in the interface.

5.5.2 Discussion

Alternatively the EPIT timer could have been used to implement this functionality. Both EPIT and GPT timers on the Sabre have a very similar interface, and both could equally have been used. However it was elected that the GPT is slightly simpler and fulfills all

requirements.

The implementation trades-off registration and removal performance for interrupt performance: registration and removal are both linear time operations as they ensure the callback timeouts are sorted based on the time they must trigger. The interrupt handler offers $O(1)$ performance for the lookup for the expiring callbacks.